

5) i) Problem Statement :

For a given set of training data examples stored in a .csv file, implement and demonstrate the find-S algorithm to output a description of the set of all hypothesis consistent.

ii) Objective :

The purpose of this program is to implement the find-S algorithm for finding the most specific hypothesis from a given set of training data stored in a .csv file.

iii) Algorithm :

* Import Necessary Libraries : Import the pandas library for handling the .csv file

* Define Find-S Algorithm Function : Implement the find-S algorithm as a function that reads data from a .csv file

c) Initialize Hypothesis : start with the most general hypothesis

d) update Hypothesis : For each positive example, refine the hypothesis by comparing attribute values

e) Return Final Hypothesis : Output the most specific hypothesis consistent with all positive training examples.

iv) Source / Program Code

```
import numpy as np
import matplotlib.pyplot as plt
from collections import Counter
```

```
data = np.random.rand(100)
```

```
labels = ["Class1" if x <= 0.5 else "Class2" for
          x in data[:50]]
```

```
def euclidean_distance(x1, x2):
    return abs(x1 - x2)
```

```
def knn_classifier(train_data, train_labels, test_point, k):
```

```
    distances = [(euclidean_distance(test_point, train_data[i]), train_labels[i]) for i in
                  range(len(train_data))]
```

```
    distance.sort(key = lambda x: x[0])
```

```
    k_nearest_neighbours = distance[:k]
```

```
    k_nearest_labels = [label for _, label in k_nearest_neighbours]
```

```
    return Counter(k_nearest_labels).most_common(1)
                           [0][0]
```

```
train_data = data[:50]
```

```
train_labels = labels
```



```
test_data = data[50:]
```

```
k_values = [1, 2, 3, 4, 5, 20, 30]
```

```
print ("--- k-Nearest Neighbours Classification ---")
```

```
print ("Training dataset: First 50 points labeled
```

```
- based on the rule ( $x \leq 0.5 \rightarrow \text{class1}$ ,  $x > 0.5$   
-  $\rightarrow \text{class2}$ )")
```

```
print ("Testing dataset: Remaining 50 points to be  
- classified \n")
```

```
results = {}
```

```
for k in k_values:
```

```
    print ("Results for k = {k}")
```

```
    classified_labels = [knn_classifier(train_data, train  
- # labels, test_point, k) for test_point  
- in test_data]
```

```
    results[k] = classified_labels
```

```
    for i, label in enumerate(classified_labels, start=
```

```
        print ("Point x[i] (value: {test_data[i]})
```

```
        - {i-51} is classified as {label}")
```

```
    print ("\n")
```

```
print ("Classification complete.\n")
```

```
for k in k_values:
```

```
    classified_labels = results[k]
```

```
    class1_points = [test_data[i] for i in range(len  
- (test_data)) if classified_labels[i] == 'class1']
```


Name of the Experiment :

```
class2-points = [test-data[i] for i in range(len
(test-data)) if classified-labels[i] == "class2"]
```

```
plt.figure(figsize=(10,6))
```

```
plt.scatter(train-data, [0] * len(train-data), c=
- ["blue" if label == "class1" else "red" for
- label in train-labels],
```

```
label="Training Data", marker="o")
```

```
plt.scatter(class1-points, [1] * len(class1-points),
- c="blue", label="class1 (Test)", marker
= "x")
```

```
plt.scatter(class2-points, [1] * len(class2-points), c
- = "red", label="class2 (Test)", marker="x")
```

```
plt.title(f"K-NN Classification Results for k={k}")
```

```
plt.xlabel("Data Points")
```

```
plt.ylabel("Classification Level")
```

```
plt.legend()
```

```
plt.grid(True)
```

```
plt.show()
```

v) Compilation and Execution steps

- * Prepare a CSV file (e.g., training-data.csv) with training data examples.
- * Save the program in a python file (e.g., find-s.py)
- * open a terminal or command prompt
- * Navigate to the directory containing the python file and CSV file
- * Run the program using the command:

python find-s.py

vi) Sample Input and Output:

Sample Input:

A .CSV file containing training data examples with various attributes and a class label (Yes/No)

Sample Output:

The final hypothesis generated by the Find-S algorithm, displayed as a list of attribute values or '?' indicating generalization

vii) Explanation of Output:

The output shows the specific hypothesis that covers all positive training examples. If a value differs between positive examples, the hypothesis will have a '?' at that position, indicating generalization.

viii) Observation and Analysis:

- * The Find-S Algorithm produces a hypothesis that is as specific as possible, covering only positive examples
- * It cannot handle noisy data or negative examples properly.

Experiment No. :

Date:

Page No.:

Name of the Experiment :

ix) Conclusion:

The find-B algorithm was successfully implemented and tested with a .CSV file, the algorithm provides the most specific hypothesis that fits all positive training examples. However, it is sensitive to noise and incomplete data.

6) i) Problem Statement :

Develop a program to implement locally weighted Regression (LWR) to fit a curve to randomly generated data points. Use a Gaussian kernel to compute the weights and visualize the resulting curve.

ii) Objective :

The purpose of this program is to implement locally weighted Regression to predict values using a weighted linear regression technique where weights decrease with distance from the query point.

iii) Algorithm :

- a) Import necessary libraries : import numpy and matplotlib for numerical computation and visualization
- b) Generate Data : Create random data points following a sinusoidal pattern with noise
- c) Define Gaussian Kernel Function : Calculate the weight for each ~~car~~ point based on its distance from the query point
- d) Implement LWR Algorithm : Compute weighted linear regression for each test point using calculated weights
- e) Visualize Results : Plot the predicted curve along with training data points

Source / Program Code

```
import numpy as np
import matplotlib.pyplot as plt
```

```
def gaussian_kernel(x, xi, tau):
    return np.exp(-np.sum((x - xi)**2) /
                    - (2 * tau ** 2))
```

```
def locally_weighted_regression(x, X, y, tau):
    m = X.shape[0]
    weights = np.array([gaussian_kernel(x, X[i],
                                         tau) for i in range(m)])
    w = np.diag(weights)
    X_transpose_W = X.T @ w
    theta = np.linalg.inv(X_transpose_W @ X) @
            X_transpose_W @ y
    return x @ theta
```

```
np.random.seed(42)
```

```
X = np.linspace(0, 2 * np.pi, 100)
```

```
y = np.sin(X) + 0.1 * np.random.randn(100)
```

```
X_bias = np.c_[np.ones(X.shape), X]
```

```
X_test = np.linspace(0, 2 * np.pi, 200)
```

```
X_test_bias = np.c_[np.ones(X_test.shape), X_test]
```

```
tau = 0.5
```

```
y_pred = np.array([locally_weighted_regression(xi,
                                                X_test_bias, y, tau) for xi in X_test_bias])
```



```
plt.figure(figsize = (10, 6))
plt.scatter(X, y, color = 'red', label = 'Training Data',
            - alpha = 0.7)
plt.plot(x-test, y-pred, color = 'blue', label = f'LWR
            - Fit (tau = {tau})', linewidth=2)
plt.xlabel('X', fontsize = 12)
plt.ylabel('y', fontsize = 12)
plt.title('Locally Weighted Regression', fontsize = 14)
plt.legend(fontsize = 10)
plt.grid(alpha = 0.3)
plt.show()
```

v) Compilation and Execution steps :

- * Save the program in a python file (e.g. lwr.py)
- * open a terminal or command prompt.
- * Navigate to the directory containing the python file
- * Run the program using the Command

```
python lwr.py
```

vi) Sample Input and output:

Sample Input:

- * 100 randomly generated data points along a sinusoidal curve
- * $\tau = 0.5$

Sample output:

- * A smooth curve fitted to the noisy data points using locally weighted Regression.

vii) Explanation of output:

The output displays a fitted curve that closely follows the underlying sinusoidal pattern, demonstrating the effectiveness of LWR for non-linear data fitting.

viii) Observation and Analysis:

- * Smaller values of τ result in a closer fit to the training data but may cause overfitting.
- * Larger values of τ produce a smoother curve but may underfit the data.

xi) Conclusion:

The Locally Weighted Regression algorithm was successfully implemented and visualized. The model effectively fits a smooth curve to the data, demonstrating the capability of LWR for non-parametric regression.